



## UNITED STATES ▾

# UNIX AS A SECOND LANGUAGE

By Sandra Henry-Stocker, Unix Dweeb, Network World

JUN 25, 2020 8:59 AM PDT

## How to stress test your Linux system

Stressing your Linux servers can be a good idea if you'd like to see how well they function when they're loaded down. In this post, we'll look at some tools that can help you add stress and gauge the results.

Why would you ever want to stress your Linux system? Because sometimes you might want to know how a system will behave when it's under a lot of pressure due to a large number of running processes, heavy network traffic, excessive memory use and so on. This kind of testing can help to ensure that a system is ready to "go public".

If you need to predict how long applications might take to respond and what, if any, processes might fail or run slowly under a heavy load, doing the stress testing up front can be a very good idea.

[Get regularly scheduled insights by signing up for Network World newsletters.]

Fortunately for those who need to be able to predict how a Linux system will react under stress, there are some helpful techniques you can employ and tools that you can use to make the process easier. In this post, we examine a few options.

## Do it yourself loops

This first technique involves running some loops on the command line and watching how they affect the system. This technique burdens the CPUs by greatly increasing the load. The results can easily be seen using the **uptime** or similar commands.

In the command below, we kick off four endless loops. You can increase the number of loops by adding digits or using a **bash** expression like **{1..6}** in place of "1 2 3 4".

```
for i in 1 2 3 4; do while : ; do : ; done & done
```

Typed on the command line, this command will start four endless loops in the background.

```
$ for i in 1 2 3 4; do while : ; do : ; done & done
[1] 205012
[2] 205013
[3] 205014
[4] 205015
```

In this case, jobs 1-4 were kicked off. Both the job numbers and process IDs are displayed.

To observe the effect on load averages, use a command like the one shown below. In this case, the **uptime** command is run every 30 seconds:

```
$ while true; do uptime; sleep 30; done
```

If you intend to run tests like this periodically, you can put the loop command into a script:

```
#!/bin/bash

while true
do
    uptime
    sleep 30
done
```


In the output, you can see how the load averages increase and then start going down again once the loops have been ended.

```
11:25:34 up 5 days, 17:27, 2 users, load average: 0.15, 0.14, 0.08
11:26:04 up 5 days, 17:27, 2 users, load average: 0.09, 0.12, 0.08
11:26:34 up 5 days, 17:28, 2 users, load average: 1.42, 0.43, 0.18
11:27:04 up 5 days, 17:28, 2 users, load average: 2.50, 0.79, 0.31
11:27:34 up 5 days, 17:29, 2 users, load average: 3.09, 1.10, 0.43
11:28:04 up 5 days, 17:29, 2 users, load average: 3.45, 1.38, 0.54
11:28:34 up 5 days, 17:30, 2 users, load average: 3.67, 1.63, 0.66
11:29:04 up 5 days, 17:30, 2 users, load average: 3.80, 1.86, 0.76
11:29:34 up 5 days, 17:31, 2 users, load average: 3.88, 2.06, 0.87
11:30:04 up 5 days, 17:31, 2 users, load average: 3.93, 2.25, 0.97
11:30:34 up 5 days, 17:32, 2 users, load average: 3.64, 2.35, 1.04 <== loops
11:31:04 up 5 days, 17:32, 2 users, load average: 2.20, 2.13, 1.01      stopped
11:31:34 up 5 days, 17:33, 2 users, load average: 1.40, 1.94, 0.98
```

Because the loads shown represent averages over 1, 5 and 15 minutes, the values will take a while to go back to what is likely normal for the system.

To stop the loops, issue a **kill** command like this one below – assuming the job numbers are 1-4 as was shown earlier in this post. If you're unsure, use the **jobs** command to verify the job IDs.

```
$ kill %1 %2 %3 %4
```





```
$ cat watch-it-2
```

UNITED STATES ▾

```
#!/bin/bash

while true
do
    free
    sleep 30
done
```

Kicking off and observing the stress:

```
$ stress --cpu 2
```

```
$ ./watch-it
13:09:14 up 5 days, 19:10, 2 users, load average: 0.00, 0.00, 0.00
13:09:44 up 5 days, 19:11, 2 users, load average: 0.68, 0.16, 0.05
13:10:14 up 5 days, 19:11, 2 users, load average: 1.20, 0.34, 0.12
13:10:44 up 5 days, 19:12, 2 users, load average: 1.52, 0.50, 0.18
13:11:14 up 5 days, 19:12, 2 users, load average: 1.71, 0.64, 0.24
13:11:44 up 5 days, 19:13, 2 users, load average: 1.83, 0.77, 0.30
```

The more CPUs specified on the command line, the faster the load will ramp up.

```
$ stress --cpu 4
$ ./watch-it
13:47:49 up 5 days, 19:49, 2 users, load average: 0.00, 0.00, 0.00
13:48:19 up 5 days, 19:49, 2 users, load average: 1.58, 0.38, 0.13
13:48:49 up 5 days, 19:50, 2 users, load average: 2.61, 0.75, 0.26
13:49:19 up 5 days, 19:50, 2 users, load average: 3.16, 1.06, 0.38
13:49:49 up 5 days, 19:51, 2 users, load average: 3.49, 1.34, 0.50
13:50:19 up 5 days, 19:51, 2 users, load average: 3.69, 1.60, 0.61
```

The **stress** command can also stress the system by adding I/O and memory load with its **--io** (input/output) and **--vm** (memory) options.

In this next example, this command for adding memory stress is run, and then the **watch-it-2** script is started:

```
$ stress --vm 2
```

```
$ watch-it-2
          UNITED STATES ▾
      total     used      free      shared  buff/cache   available
Mem:  6087064  662160  2519164      8868    2905740  5117548
Swap: 2097148      0  2097148
      total     used      free      shared  buff/cache   available
Mem:  6087064  803464  2377832      8864    2905768  4976248
Swap: 2097148      0  2097148
      total     used      free      shared  buff/cache   available
Mem:  6087064  968512  2212772      8864    2905780  4811200
Swap: 2097148      0  2097148
```

Another option for **stress** is to use the **--io** option to add input/output activity to the system. In this case, you would use a command like this:

```
$ stress --io 4
```

You could then observe the stressed IO using **iostop**. Note that **iostop** requires root privilege.

### before

```
$ sudo iostop -o
Total DISK READ:      0.00 B/s | Total DISK WRITE:      19.36 K/s
Current DISK READ:     0.00 B/s | Current DISK WRITE:     27.10 K/s
      TID  PRIO  USER      DISK READ  DISK WRITE  SWAPIN     IO>     COMMAND
269308  be/4  root      0.00 B/s     0.00 B/s    0.00 %  1.24 % [kworker~fficient]
      283  be/3  root      0.00 B/s    19.36 K/s    0.00 %  0.26 % [jbd2/sda1-8]
```

### after

```
Total DISK READ:      0.00 B/s | Total DISK WRITE:      0.00 B/s
Current DISK READ:     0.00 B/s | Current DISK WRITE:     0.00 B/s
      TID  PRIO  USER      DISK READ  DISK WRITE  SWAPIN     IO>     COMMAND
270983  be/4  shs      0.00 B/s     0.00 B/s    0.00 %  51.45 % stress --io 4
270984  be/4  shs      0.00 B/s     0.00 B/s    0.00 %  51.36 % stress --io 4
270985  be/4  shs      0.00 B/s     0.00 B/s    0.00 %  50.95 % stress --io 4
270982  be/4  shs      0.00 B/s     0.00 B/s    0.00 %  50.80 % stress --io 4
269308  be/4  root     0.00 B/s     0.00 B/s    0.00 %  0.09 % [kworker~fficient]
```

**Stress** is just one of a number of tools for adding stress to a system. Another and newer tool, **stress-ng**, will be covered in a future post.

## Wrap-Up

Various tools for stress-testing a system will help you anticipate how systems will respond in real world situations in which they are subjected to increased traffic and computing demands.

While what we've shown in the post are ways to create and measure various types of stress, the ultimate benefit is how the stress helps in determining how well your system or application responds to it.

UNITED STATES ▾  
Join the Network World communities on [Facebook](#) and [LinkedIn](#) to comment on topics that are top of mind.

Sandra Henry-Stocker has been administering Unix systems for more than 30 years. She describes herself as "USL" (Unix as a second language) but remembers enough English to write books and buy groceries. She lives in the mountains in Virginia where, when not working with or writing about Unix, she's chasing the bears away from her bird feeders.

Follow

Copyright © 2020 IDG Communications, Inc.

## 💡 SD-WAN buyers guide: Key questions to ask vendors (and yourself)

### SPONSORED LINKS

[Automate IT tasks and Free Up Your Time with Jamf Now](#)

[dtSearch® instantly searches terabytes of files, emails, databases, web data. See site for hundreds of reviews; enterprise & developer evaluations](#)

[Getting a Grip on Basic Cyber Hygiene with the CIS Controls](#)

[Your cloud, your way: Why Cloud Verified matters](#)

[Don't Pay the Ransom! Learn how to fight back against cyber criminals.](#)

[Solution Architect Perspectives: Cybersecurity in a Perimeter-less Network Environment](#)



Copyright © 2021 IDG Communications, Inc.